

MARK BASOV

basovmark@gmail.com ♦ [LinkedIn](#) ♦ [Telegram](#)

OBJECTIVE

ML Engineer with **4 years** of commercial experience executing ML and analytics projects in NLP and classic ML. Primarily works with LLMs: SFT, DPO, KTO, QLoRA, inference optimization, prompt enhancing. Has deep expertise in RAG. Worked both remotely and hybrid in teams of 2 to 16 people. Convert unclear business requirements into production services. Handle R&D work – can quickly pivot between different development directions. Actively explore cutting-edge technologies. Motivated by delivering business value and consider work an essential part of my identity.

SKILLS

ML & Analytics	NumPy, pandas, Polars, scikit-learn, CatBoost, LightGBM, statsmodels, matplotlib, seaborn, Plotly
LLM & NLP	PyTorch, Transformers, LangChain, spaCy, pymorphy3, vLLM, SGLang, Haystack, PEFT, trl
Other	Python, PostgreSQL, ClickHouse, Kubeflow, Oneflow, Tensorboard, PySpark, FastAPI, Pydantic, FAISS, Pinecone, Git, Docker, k8s, Linux, Bash, DVC, ClearML, MLflow, unittest, ms-swift, Jira
Languages	Russian (native), English (B2)

EXPERIENCE

VK Senior ML Engineer	Sep 2025 - Present <i>Remote</i>
---------------------------------	-------------------------------------

- Saved **>5M RUB** and reduced labeling time **30×** by replacing human assessors with an on-prem ML judge. Achieved **97%** agreement with human labels, enabling scalable data labeling without external API costs.
- Reduced no-context answers by **70%** and lifted retrieval **recall@5 by 20%** by training a cross-encoder reranker and integrating an LLM-based chunk filtering layer into the production RAG pipeline.
- Increased **resolution rate by 13%** for answered queries by aligning the generative model with RAG-DDR; eliminated human-in-the-loop labeling and enabled **fully automated fine-tuning** for new customers.
- Standardized document processing for **8** business units by migrating to a unified Docling-based pipeline; added robust support for PDF and Markdown parsing.

Stack: Python, PyTorch, Peft, trl, vLLM, k8s, SQL, Kubeflow, Transformers, ms-swift, DVC, Docling, Oneflow

2GIS ML Engineer	May 2024 - Sep 2025 <i>Remote</i>
----------------------------	--------------------------------------

- Increased card openings by **12%** and conversion to actions by **8%** through developing ML attribute matching system. Built production service for attribute processing, scaled to 50 regions and 4 domains.
- Reduced category analysis time from 6-12 months to **7-14 days** and decreased error rate by **15%** through ML pipeline for fact extraction from multimodal content. Built from scratch an analytics platform.
- Developed functional MVPs as part of R&D:
 - Reduced category development time by **20%** by creating keyword generator from search queries.
 - Increased overall product parsing completeness by **5%** by fine-tuning NER parser for HTML pages.
 - Identified **6%** of organizations with incorrect working hours based on banking transactions.
 - Implemented LLM segmentor for search queries into structured API filters.

Stack: Python, LangChain, PyTorch, Transformers, pandas, scikit-learn, k8s, Docker, PostgreSQL, PEFT, ClearML

Softline Data Scientist	Mar 2022 - May 2024 <i>Moscow, Russia</i>
-----------------------------------	--

- Reduced expert manual effort by **40+** hours monthly by developing an ML-based recommendation system. Reduced waste on unnecessary communications by **4x**.
- Improved the personal marketing communications strategy of a major Russian retailer. The treatment group turned out to be **30%** more profitable than the control group.
- Implemented MLOps pipeline for credit scoring reducing deployment cycle **from weeks to hours**.

Stack: Python, NumPy, pandas, CatBoost, LightGBM, scikit-learn, MLflow, statsmodels, matplotlib, Plotly, PostgreSQL

EDUCATION

ITMO University, Remote

2024 - 2026

Master's degree in Artificial Intelligence.

ITMO University, Saint-Petersburg

2020 - 2024

Bachelor's degree in Cybersecurity. Honors diploma.